


Manual Detector-Finding Review: Rating Distribution

263 rated findings across 4 judge models (DeepSeek, OpenAI, OpenAI-5.6, Qwen)



Model	n	Accurate	Likely accurate	Partially accurate	Inaccurate
Overall	263	84	148	12	19
DeepSeek	56	10	32	0	14
OpenAI	110	27	74	9	0
OpenAI-5.6	62	31	25	1	5
Qwen	35	16	17	2	0

Notes: DeepSeek shows the highest inaccurate rate ($14/56 = 25\%$) and zero partially-accurate findings. OpenAI has zero inaccurate findings but the most partially-accurate ($9/110 = 8\%$). OpenAI-5.6 and Qwen are both mostly clean, with OpenAI-5.6 showing 5 inaccurate (8%) and Qwen only 2 partially-accurate and 0 inaccurate. Overall, ~88% of findings were rated accurate or likely accurate.